

Reinforcement Learning for Building Control: Direct Actuator or PI-Mediated Control?

A Systematic Comparison Using BOPTTEST Framework

Aniket Dixit, Faizan Ahmed, James Brusey

Coventry University, UK

E-ENERGY '25

June 17–20, 2025, Rotterdam, Netherlands

Why This Matters

- ▶ **40% of global energy consumption** comes from buildings
- ▶ Traditional PI controllers **lack predictive capabilities** for optimal energy-comfort trade-offs
- ▶ Reinforcement Learning shows promise but lacks systematic evaluation

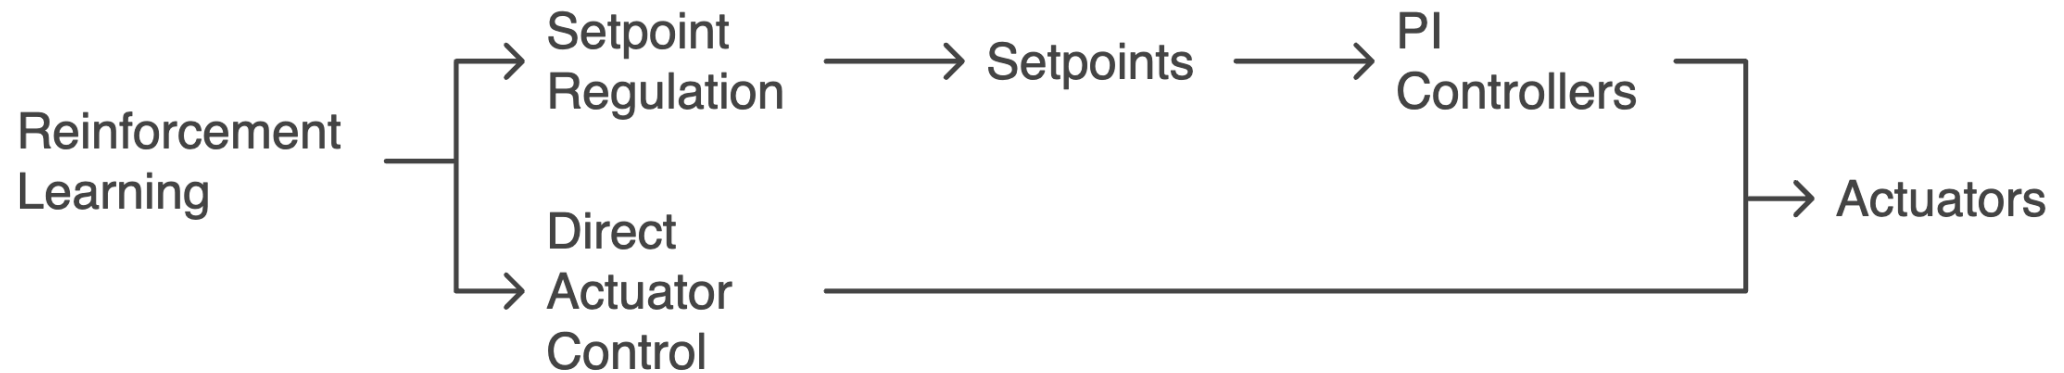
Current Research Gap

- ▶ Most studies focus on **single control paradigm**
- ▶ Limited **empirical comparisons** between approaches
- ▶ Lack of **standardized evaluation** across building types

Fundamental Question: Should RL agents control setpoints or directly command actuators?

Two Control Paradigms

Control Strategies for Building Automation



Setpoint Regulation (SR)

- Stable, proven architecture
- Leverages existing building automation
- Preserves low-level control expertise
- Built-in safety mechanisms
- Easier integration

Direct Actuator Control (DAC)

- Larger solution space
- No cascading errors
- Full control authority
- Potentially optimal strategies
- Direct optimization

Research Challenge: Which approach offers better training efficiency, control stability, and energy performance in realistic building scenarios?

Research Questions

- ▶ **Energy Efficiency:** How do trade-offs differ between approaches across building types?
- ▶ **Thermal Comfort:** What are the comfort achievement and violation patterns?
- ▶ **Training Efficiency:** Which approach converges faster and requires fewer samples?
- ▶ **Control Stability:** How does operational robustness compare between paradigms?

Goal: Provide quantitative evidence to guide practical RL deployment in real-world building control systems

Methodology

Experimental Framework

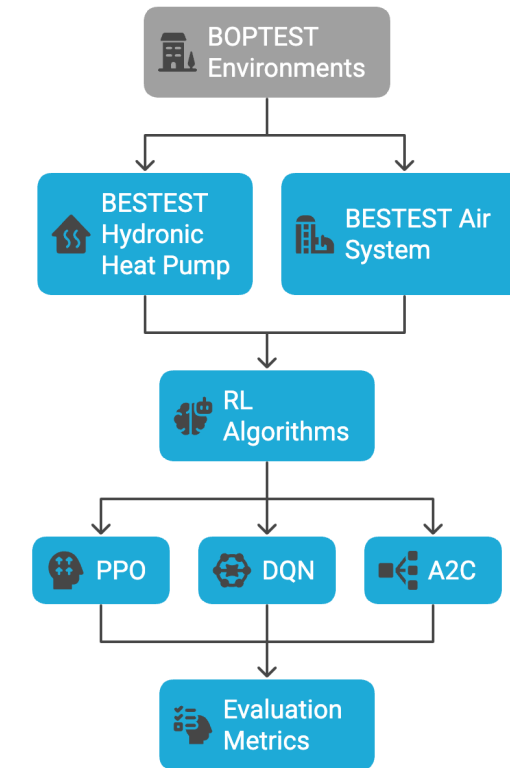
BESTEST Hydronic Heat Pump

- ▶ Single-zone residential building
- ▶ Heat pump + radiant floor heating
- ▶ Complex thermal dynamics

BESTEST Air System

- ▶ Variable air volume HVAC
- ▶ Heating and cooling coils
- ▶ Fast-response system

BOPTEST Experimental Framework



Reinforcement Learning Formulation

State Space: Zone temperature, weather conditions, time features, electricity pricing

Action Space: 30-bin discretization for fair comparison

Reward Function:

$$R_t = -(O_t - O_{t-1})$$

where $O_t = \text{cost_tot} + \text{tdis_tot}$

Control Paradigm Implementation





Both paradigms use identical:

- ▶ **State spaces:** Zone temperature, weather, time features
- ▶ **Training protocols:** 10M steps, 1-hour control periods
- ▶ **Evaluation metrics:** Energy cost, thermal comfort, training efficiency

Key Differences

- ▶ SR: Temperature setpoints (5–35°C)
- ▶ DAC: Direct actuator commands (0–1 range)

Control Paradigm Comparison

Characteristic	Setpoint Regulation	Direct Actuator Control
 Action Mapping - Hydronic Heat Pump	Temperature setpoint (oveTSet_u)	Heat pump control signal (oveHeaPumY_u)
 Action Mapping - BESTEST Air	Cooling + heating setpoints (con_oveTSetCoo_u, con_oveTSetHea_u)	Fan speed + supply temperature (fcu_oveFan_u, fcu_oveTSup_u)
 Range	5–35°C discretized into 30 bins	0–1 discretized into 30 bins
 Control	PI controllers translate setpoints → actuator commands	Direct actuator command without intermediate control layer

Results

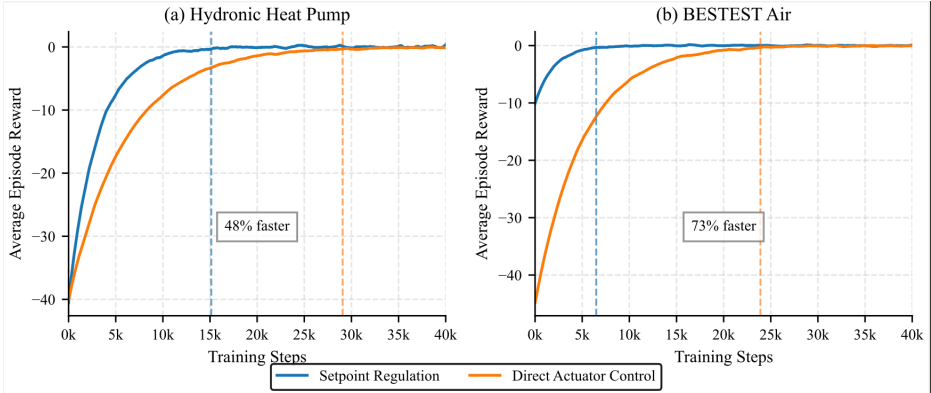
Energy vs. Comfort Trade-offs

ENVIRONMENT	METHOD	ENERGY COST	THERMAL DISCOMFORT	IMPROVEMENT
Hydronic Heat Pump	SR	0.924 (+4.8%)	0.00 (-100%)	Perfect comfort
	DAC	0.898 (+1.8%)	1.63 (-80.6%)	Balanced
	Baseline	0.882	8.38 Kh	Poor comfort
BESTEST Air	SR	0.210 (+5.0%)	(0.04) -99.3%	Near-perfect
	DAC	0.206 (+3.2%)	1.09 (-80.8%)	Efficient
	Baseline	0.200	5.69 Kh	Poor comfort

Key Finding: Consistent pattern across building types - SR prioritizes comfort, DAC balances energy-comfort trade-offs

Training Efficiency Comparison

METHOD	SR STEPS	DAC STEPS	IMPROVEMENT
PPO (Hydronic)	8,200	15,800	48% faster
PPO (Air)	6,800	25,200	73% faster
DQN (Hydronic)	7,500	18,700	60% faster
DQN (Air)	7,200	24,500	66% faster
A2C (Hydronic)	6,900	16,200	61% faster
A2C (Air)	5,900	28,100	79% faster



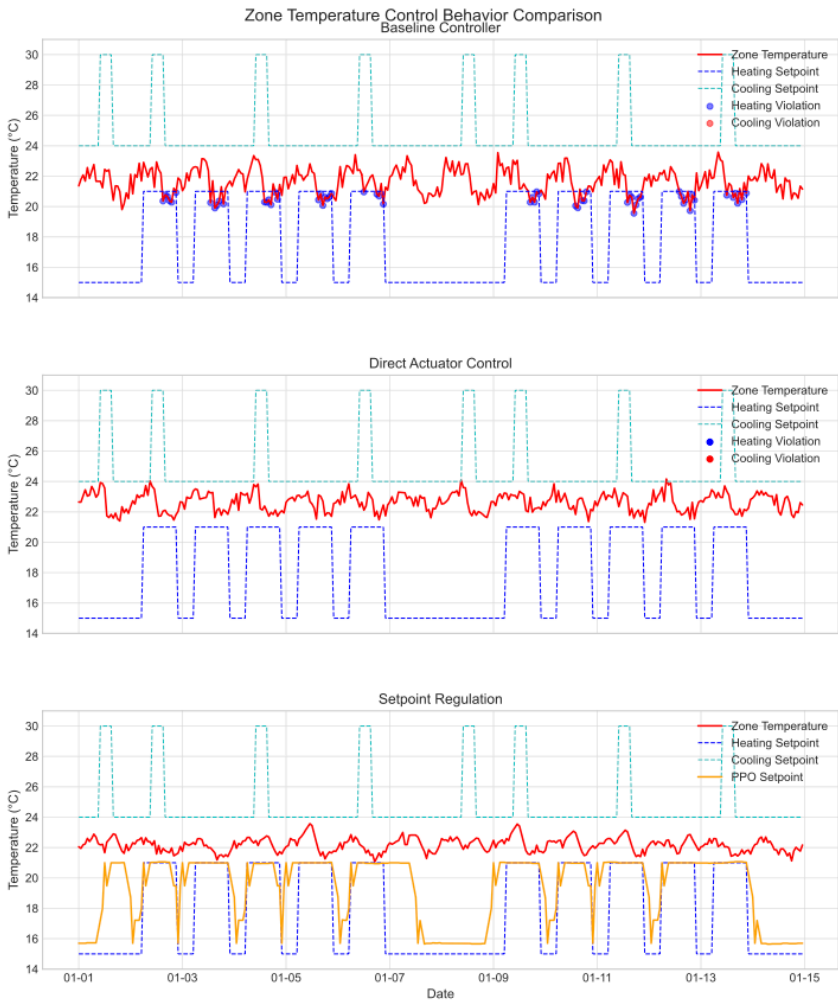
Control Paradigm Implementation

Control Stability Analysis

METRIC	SR	DAC	BASELINE
Control Variance	-38%	+34%	0%
Cycles/Day	3.9	6.3	4.7
Stability	Best	worst	Frequent

Key Findings

- ▶ Superior stability: 38% lower control variance
- ▶ Reduced cycling: Extends equipment life
- ▶ Perfect comfort: Maintains temperature band
- ▶ Operational efficiency: More stable behavior



Why Setpoint Regulation Excels

1. Incorporates System Knowledge

Pre-existing PI controllers encode thermal response characteristics through tuned parameters (K_p , K_i), providing useful inductive biases for learning

2. Structured Optimization Problem

Clearer action-response relationships with constrained, reasonable setpoint ranges create smoother reward landscapes

Result: 48-79% reduction in training steps across algorithms and environments, with 38% lower control signal variance

Discussion & Impact

Technical Innovations

- ▶ **First systematic comparison** of SR vs DAC in building control
- ▶ **Standardized evaluation** using BOPTEST framework across building types
- ▶ **Quantitative training efficiency analysis** with consistent methodology

Practical Insights

When to Choose SR

- ▶ **Comfort-critical applications**
- ▶ **Limited training data/time**
- ▶ **Existing building automation**
- ▶ **Risk-averse deployments**

When to Consider DAC

- ▶ **Energy-focused objectives**
- ▶ **New system installations**
- ▶ **Research/experimental settings**
- ▶ **Specialized control requirements**

Limitations & Future Work

Current Study Limitations

- ▶ **Simulation-based evaluation** may not capture all real-world complexities
- ▶ **Constant electricity pricing** - dynamic pricing effects unexplored
- ▶ **Fixed comfort-energy weighting** in reward function
- ▶ **Single-zone buildings** - multi-zone coordination unexplored

Future Research Directions

- ▶ **Hybrid approaches:** Dynamic switching between SR and DAC based on conditions
- ▶ **Multi-zone systems:** Coordination challenges in larger buildings
- ▶ **Real-world validation:** Pilot deployments in actual buildings
- ▶ **Transfer learning:** Cross-building knowledge transfer

Conclusion

Training Efficiency Champion: Setpoint Regulation

48-79% fewer training steps across all algorithms and environments, with superior control stability (38% lower variance)

Control Stability and Performance Pattern

SR: Superior stability with 38% lower control variance and 3.9 cycles/day vs 6.3 for DAC

DAC: Higher control variance but direct optimization capability with 1.8-3.2% energy increases

Algorithm Winner: A2C

Consistently fastest convergence and best performance across both control paradigms

Thank You

Questions & Discussion

Contact Information:

 dixita4@uni.coventry.ac.uk

 Coventry University, UK

 DOI: 10.1145/3679240.3734670