

# Smart HVAC Systems: RL Approach

**Aniket Dixit**

PhD Candidate, Computer Science

Tech University

June 15, 2025

**Advisor:** Prof. James Brusey

**Committee:** Prof. James Brusey, Prof. Matthew England

# Research Overview

- ▶ **Problem:** Current sequential models struggle with long-range dependencies
- ▶ **Approach:** Novel attention mechanism with adaptive memory
- ▶ **Contribution:** 15% improvement over transformer baselines
- ▶ **Impact:** Applications in machine translation, summarization, and dialogue
- ▶ **Validation:** Comprehensive evaluation across 8 benchmark datasets

**Core Innovation:** Dynamic attention weights that adapt based on sequence complexity and context relevance

# Problem Statement

# Sequential Data Processing Challenges

## Current Limitations

- ▶ **Computational Complexity:** Standard attention scales quadratically with sequence length
- ▶ **Memory Constraints:** Limited ability to maintain long-term context
- ▶ **Training Instability:** Gradient vanishing in very deep architectures
- ▶ **Domain Adaptation:** Poor generalization across different text domains

## Research Questions

- ▶ How can we design attention mechanisms that scale efficiently?
- ▶ What architectural changes improve long-range dependency modeling?
- ▶ Can we achieve better performance with fewer parameters?

**Hypothesis:** Adaptive attention with hierarchical memory structures can overcome current limitations while maintaining computational efficiency

# Reward Function

We model sequential data processing as learning a mapping function:

**Objective Function:**

$$\mathcal{L}(\theta) = \sum_{i=1}^N \ell(f_{\theta}(X^{(i)}), Y^{(i)}) + \lambda \Omega(\theta)$$

Where  $f_{\theta}$  represents our proposed attention-based architecture,  $\ell$  is the loss function, and  $\Omega(\theta)$  is a regularization term.

# Methodology

# Proposed Architecture: AdaptiveNet

## Key Components

- ▶ **Multi-Head Adaptive Attention:** Dynamic attention weights
- ▶ **Hierarchical Memory Module:** Long-term context storage
- ▶ **Residual Gating:** Improved gradient flow
- ▶ **Layer-wise Learning Rates:** Optimized training dynamics

**\*\*AdaptiveNet Architecture\*\***  
[Insert detailed network diagram]



## Innovation Details

The adaptive attention mechanism computes context-dependent weights:

$$\alpha_{ij} = \frac{\exp(e_{ij} \cdot \text{adapt}(c_i))}{\sum_{k=1}^n \exp(e_{ik} \cdot \text{adapt}(c_i))}$$

# Training Algorithm

## Algorithm 1: AdaptiveNet Training

**Input:** Training dataset  $\mathcal{D}$ , model parameters  $\theta_0$ , learning rate  $\eta$

**Output:** Optimized parameters  $\theta^*$

1. **Initialize** model parameters  $\theta_0$  and memory states  $M_0$
2. **for** epoch = 1 to max\_epochs **do**
3.    $\text{AdaptiveNet}(X, M)$   
    and pass:  $\hat{Y} = \text{AdaptiveNet}(X, A,$
4. **end for**

# Theoretical Analysis

## Theorem 1: Convergence Guarantee

Under standard assumptions of bounded gradients and Lipschitz continuity, the AdaptiveNet training algorithm converges to a stationary point with probability 1.

### Proof Sketch:

1. Show that the adaptive attention preserves the contraction property
2. Apply stochastic approximation theory for the memory update rule
3. Use martingale convergence theorem for the parameter updates

## Complexity Analysis

- ▶ **Time Complexity:**  $O(n \log n)$  for sequence length  $n$  (vs.  $O(n^2)$  for standard attention)
- ▶ **Space Complexity:**  $O(n + m)$  where  $m$  is memory size
- ▶ **Parameter Efficiency:** 23% fewer parameters than comparable transformer models

# Experimental Results

# Experimental Setup

## Datasets

- ▶ **Machine Translation:** WMT14 En-De, En-Fr
- ▶ **Text Summarization:** CNN/DailyMail, XSum
- ▶ **Question Answering:** SQuAD 2.0, Natural Questions
- ▶ **Dialogue:** PersonaChat, MultiWOZ

# Baseline Models

- ▶ **Transformer-Base:** Standard attention mechanism
- ▶ **Linformer:** Linear attention approximation
- ▶ **Performer:** Fast attention via random features
- ▶ **Longformer:** Sparse attention patterns

**\*\*Experimental Pipeline\*\*** [Insert flowchart showing data preprocessing, model training, evaluation metrics, and statistical testing procedures]

# Main Results

**\*\*Performance Comparison Across Tasks\*\***  
[Figure ]

## Key Findings

TASK	ADAPTIVENET	TRANSFORMER	IMPROVEMENT
Translation (BLEU)	34.2	29.8	+14.8%
Summarization (ROUGE-L)	42.1	38.7	+8.8%
QA (F1)	89.3	85.2	+4.8%
Dialogue (BLEU)	28.6	24.1	+18.7%



# Ablation Study

## Component Analysis

COMPONENT	PERFORMANCE IMPACT	STATISTICAL SIGNIFICANCE
Adaptive Attention	+12.3%	$p < 0.001$
Hierarchical Memory	+8.7%	$p < 0.001$
Residual Gating	+4.2%	$p < 0.01$
Layer-wise LR	+2.8%	$p < 0.05$

**Critical Finding:** The adaptive attention mechanism provides the largest performance gain, with hierarchical memory being the second most important component

## Computational Efficiency

**\*\*Speed and Memory Comparison\*\*** [Insert line graphs showing training time, inference speed, and memory usage across different sequence lengths for all baseline models]

# Qualitative Analysis

## Attention Visualization

**\*\*Attention Patterns\*\*** [Insert heatmaps showing attention weights for sample sentences, demonstrating long-range dependencies and adaptive behavior]

## Example Outputs

### Translation Quality:

- ▶ **Our Model:** "This is a very complex scientific article"
- ▶ **Baseline:** "This is a complex scientific paper"

### Key Improvements:

- ▶ Better preservation of semantic meaning
- ▶ More accurate handling of technical terms
- ▶ Improved coherence in long documents

# Error Analysis

## Failure Cases

- ▶ **Very Short Sequences** (<10 tokens): Adaptive mechanism adds unnecessary overhead
- ▶ **Highly Repetitive Text**: Memory module can get stuck in local patterns
- ▶ **Code-Switching**: Limited training data for multilingual scenarios

## Robustness Testing

PERTURBATION TYPE	PERFORMANCE DROP	RECOVERY METHOD
Noise Injection	-8.2%	Data augmentation
Domain Shift	-12.5%	Few-shot adaptation
Adversarial	-15.3%	Adversarial training

# Contributions & Impact

# Primary Contributions

## Technical Innovations

- ▶ **Novel Architecture:** First adaptive attention mechanism with hierarchical memory
- ▶ **Theoretical Framework:** Convergence guarantees for adaptive training dynamics

## Scientific Impact

- ▶ **Performance:** State-of-the-art results on 6 out of 8 benchmark datasets
- ▶ **Efficiency:** 40% faster training, 60% faster inference than baselines

**Broader Impact:** This work enables deployment of advanced NLP models in resource-constrained environments while achieving superior performance

# Future Directions

## Short-term Goals (6-12 months)

- ▶ **Multimodal Extension:** Adapt architecture for vision-language tasks
- ▶ **Few-shot Learning:** Improve performance with limited training data
- ▶ **Hardware Optimization:** Develop GPU-optimized implementations

## Long-term Vision (2-5 years)

- ▶ **Foundation Models:** Scale to billion-parameter models
- ▶ **Real-time Applications:** Deploy in production systems
- ▶ **Cross-lingual Transfer:** Universal language understanding

# Collaboration Opportunities

- ▶ **Industry Partnerships:** Google, Microsoft, Meta AI research teams
- ▶ **Academic Networks:** Stanford HAI, MIT CSAIL, CMU LTI
- ▶ **Open Science:** Contributing to Hugging Face, PyTorch ecosystem



# References

1. Vaswani, A., et al. (2017). Attention is All You Need. *Neural Information Processing Systems*, 5998-6008.
2. Devlin, J., et al. (2019). BERT: Pre-training of Deep Bidirectional Transformers. *NAACL-HLT*, 4171-4186.
3. Wang, S., et al. (2020). Linformer: Self-Attention with Linear Complexity. *arXiv preprint arXiv:2006.04768*.
4. Choromanski, K., et al. (2021). Rethinking Attention with Performers. *ICLR*.
5. Beltagy, I., et al. (2020). Longformer: The Long-Document Transformer. *arXiv preprint arXiv:2004.05150*.
6. Brown, T., et al. (2020). Language Models are Few-Shot Learners. *Neural Information Processing Systems*, 1877-1901.

# Thank You

## Questions & Discussion

### Contact Information:

 [jane.smith@techuniversity.edu](mailto:jane.smith@techuniversity.edu)

 @janesmith\_ai

 [janesmith-research.com](http://janesmith-research.com)

### Resources:

 Code: [github.com/janesmith/adaptivenet](https://github.com/janesmith/adaptivenet)

 Paper: [arxiv.org/abs/2025.54321](https://arxiv.org/abs/2025.54321)

 Demo: [adaptivenet-demo.com](http://adaptivenet-demo.com)

# Appendix: Implementation Details

## Software Architecture

### Core Dependencies

- **Framework:** PyTorch 2.0 + Transformers 4.28
- **Optimization:** AdamW with cosine scheduling
- **Distributed Training:** DeepSpeed ZeRO Stage 2

### Hardware Configuration

- **Training:** 8× NVIDIA A100 (80GB) GPUs
- **Evaluation:** Single V100 (32GB) GPU
- **Storage:** High-speed NVMe SSD array

### Hyperparameter Configuration:

Learning Rate:  $1e-4$  (adaptive attention),  $5e-5$  (other components)

Batch Size: 32 per GPU (256 total)

Sequence Length: 512 (training), 1024 (evaluation)

Memory Size: 256 slots

Attention Heads: 16

Hidden Dimension: 768

Dropout: 0.1

Weight Decay: 0.01

# Appendix: Extended Results

**\*\*Detailed Performance Analysis\*\*** [Insert comprehensive results table with confidence intervals, statistical significance tests, and cross-validation scores for all experiments]

# Statistical Analysis

- ▶ **Sample Sizes:** 10 random seeds × 3 dataset splits × 5 cross-validation folds
- ▶ **Significance Testing:** Paired t-tests with Bonferroni correction
- ▶ **Effect Sizes:** Cohen's d > 0.8 for all major improvements
- ▶ **Confidence Intervals:** 95% bootstrapped intervals reported

# Computational Profiling

OPERATION	TIME (MS)	MEMORY (MB)	GPU UTILIZATION (%)
Forward Pass	12.3	2,840	87%
Backward Pass	18.7	3,120	92%
Memory Update	3.2	480	45%
Attention Computation	8.9	1,760	78%

# Thank You

## Questions & Discussion

### Contact Information:

 [your.email@university.edu]

 [@your\_twitter\_handle]

 [your-website.com]

### Resources:

 Code: [github.com/username/repo]

 Paper: [arxiv.org/abs/xxxx.xxxxx]

 Demo: [project-website.com]

